

BIBA: Business Intelligence and Big Data

2018-09-19

Jens Ulrik Hansen

Data

Questions

- Are you comfortable with your business case hand-ins and your groups?
- Are you comfortable with R and the Azure notebook?

Today's program

- Reflections on the business case hand-ins
- Types of data and tidy data
- Data transformation/ETL
- Exercises in R
- Data cleaning
- Exercises in R

Reflections on the business case hand-ins

Types of data and tidy data

The data science/BI process recap

- Using data (and data analysis) to solve business problems
 1. identify business problem
 2. collect data
 3. prepare data
 4. analyze data
 5. conclude and communicate

Data is everywhere

- In books and paper record
- Collected through surveys
- Collected through measurement and sensors (weather data etc)
- Collected on the web
- In business IT-systems and databases (ERP and CRM systems, etc.)
- Geo-location data
- Social media data
- In pictures and sounds
- In our brains and genes
- ...

Big Data

- The Big Data revolution have challenged the data landscape in several ways (-the 3 Vs of Big Data):
 - There are so much more data available than ever before (Volume)
 - The data arrive real-time, it needs to be streamed instead of just extracted once (Velocity)
 - The data comes from many different sources and in many different formats (Variety)
- This, of course, also challenge how we collect, store, extract, transform, and load data
- More on these challenges later in the course...

Data representations

- Data can be represented in a variety of ways and in a variety of formats such as
 - tables/spreadsheets, databases/SQL, No-SQL, plain text, XML, JSON, graphs, documents, ... etc.
- We will mainly work with data in a very particular format, stored in **rows** and **columns** (like a traditional spreadsheet/table)
- When rows represent ***cases/observations/objects*** and columns represent ***attributes/variables/features*** we will call the data “***tidy data***” (- terminology by Hadley Wickham - his paper “Tidy data” on Moodle)
- Examples of cases:
 - A persons
 - a sensor measurement
 - a transaction
- Examples of attributes:
 - eye color of a person
 - temperature of a sensor at particular time
 - the costumer of the transaction

Tidy data examples

- The same data can be represented as tidy in different ways
 - Student records, with students as observations
 - Classes as observations
 - Or each student enrollment in a class an observation
- Another example:

```
In [11]: mtcars
```

| | mpg | cyl | disp | hp | drat | wt | qsec | vs | am | gear | carb |
|--------------------------|------|-----|-------|-----|------|-------|-------|----|----|------|------|
| Mazda RX4 | 21.0 | 6 | 160.0 | 110 | 3.90 | 2.620 | 16.46 | 0 | 1 | 4 | 4 |
| Mazda RX4 Wag | 21.0 | 6 | 160.0 | 110 | 3.90 | 2.875 | 17.02 | 0 | 1 | 4 | 4 |
| Datsun 710 | 22.8 | 4 | 108.0 | 93 | 3.85 | 2.320 | 18.61 | 1 | 1 | 4 | 1 |
| Hornet 4 Drive | 21.4 | 6 | 258.0 | 110 | 3.08 | 3.215 | 19.44 | 1 | 0 | 3 | 1 |
| Hornet Sportabout | 18.7 | 8 | 360.0 | 175 | 3.15 | 3.440 | 17.02 | 0 | 0 | 3 | 2 |
| Valiant | 18.1 | 6 | 225.0 | 105 | 2.76 | 3.460 | 20.22 | 1 | 0 | 3 | 1 |

Types of Attributes

- There are different types of attributes
 - ***Nominal***
 - Examples: ID numbers, eye color, zip codes
 - ***Ordinal***
 - Examples: rankings (e.g., taste of potato chips on a scale from 1 to 10), grades, height in {tall, medium, short}
 - ***Interval***
 - Examples: calendar dates, temperatures in Celsius or Fahrenheit.
 - ***Ratio***
 - Examples: temperature in Kelvin, length, time, counts

Tidy data

- *“Like families, tidy datasets are all alike but every messy dataset is messy in its own way. Tidy datasets provide a standardized way to link the structure of a dataset (its physical layout) with its semantics (its meaning)”* - Hadley Wickham, “Tidy Data”.
- Every cell represent one piece of information
- Every column have same number of entries
- Each observation contains all values measured on that same unit/individual across attributes
 - Variables are what is measured in a study for each subject (observation)
- It is not always obvious what are observations and what are variables
 - However, a general rule of thumb: Easier to describe functional relationships between variables. Easier to make comparison between groups of observations.

Tidy data

- Reasons for working with tidy data
 - It provide consistency in data representation, which allow for a general set of operations on the data
 - Many natural operations becomes simple to perform (especially in R), such as constructing composite variables
- Reasons for working with non-tidy data
 - there can be reason to have data in other formats than in this tidy format. For instance, tidy data format is not always the most space efficient format. Moreover, particular computations might be performed faster on data in other formats than the tidy format. (See the following blog post by Jeff Leek for more on non-tidy data: <http://simplystatistics.org/2016/02/17/non-tidy-data/>)

Tidy data in R

- In R, data frames are exactly such row/column representations of data
- Technically, a data frame is a names list (of columns) where each list element is a vector of the same length

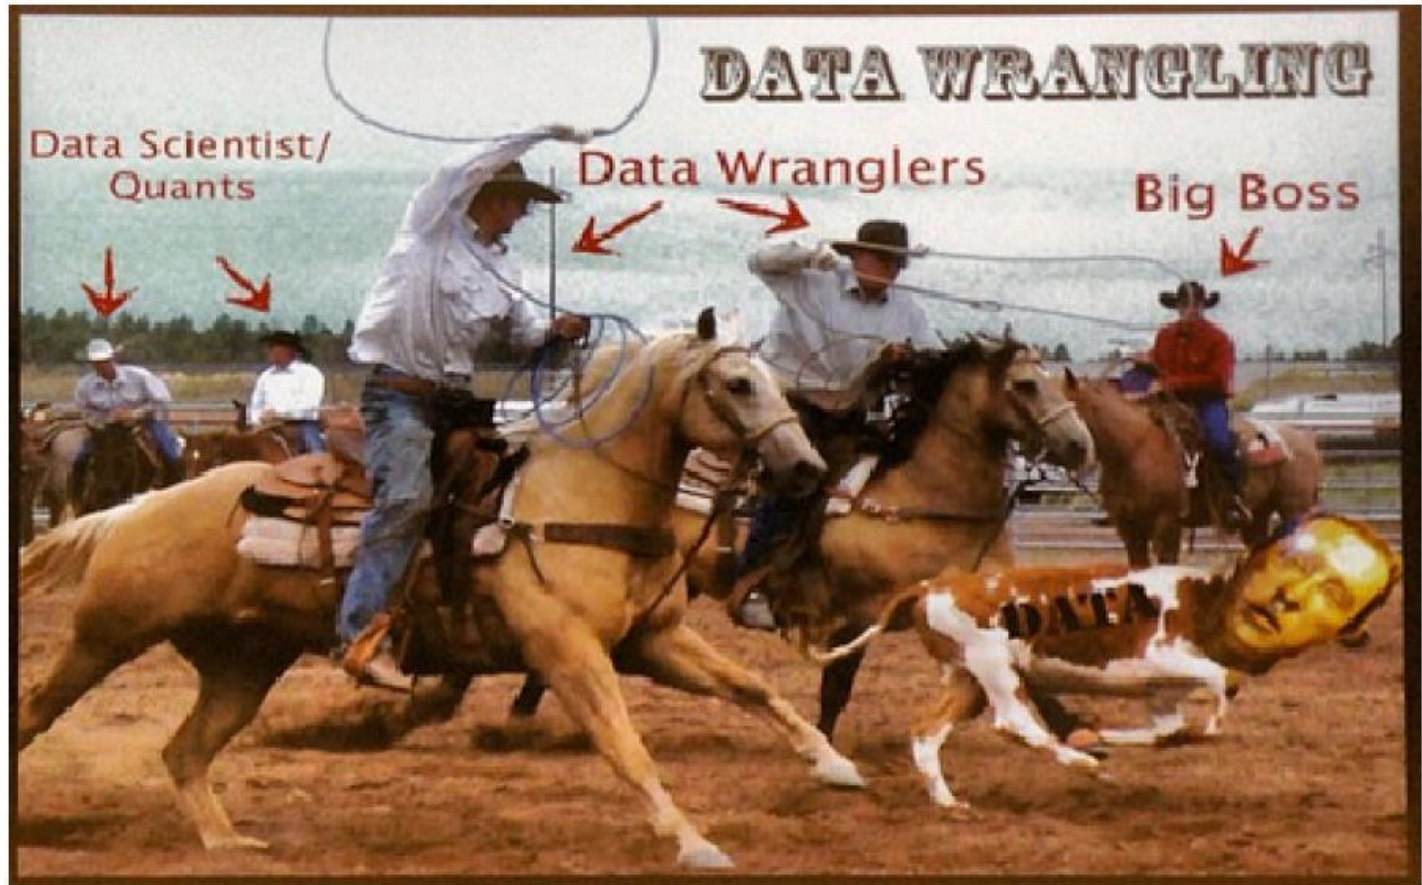
Different data types in R

- The different types of attributes can in R be represented by different data types:
 - **Numeric** (floating point numbers or integers)
 - 5.34, pi, 3.3333, 1, 5.0, 5L
 - **Categorical / factors**
 - “yes”/”no”, “Country of origin” (numbers underneath)
 - Can be given an order
 - **Dates**
 - “2016-09-22” (from the Sys.Date() command)
 - (can internally be stored as integers or floating point numbers)
 - **Time-Date**
 - "2016-09-22 12:06:30 CEST" (from the Sys.time() command)
 - **Unstructured character strings**
 - “In this part of the book, you’ll learn about data wrangling, the art of”

Data transformation/ETL

Data transformation

- Data transformation / data wrangling
- ETL – extract, transform, load
- This is often the most time consuming



Extract: Getting data into R

- Reading local files
 - CSV files
 - Excel files
 - JSON files
 - XML
- Reading from the web
 - We can use the same files
- Reading from relational databases
 - SQL
 - No-SQL
- From APIs
 - Twitter
 - Quandl
- Other ways of reading data

Reading data into R from local files

- Make sure the file is in the working directory
 - See the current one: `getwd()`
 - Set it: `setwd("../")`
- To read “.csv” files use `read.csv`
 - Note, “csv” stands for “comma separated values”. Due to the use of “,” as decimal indicator in Danish, you can run into issues.
 - `read.csv2` is just as `read.csv`, except that it is for files using semi-colon “;” for separation instead of “,”

Reading data into R

- Loading data into R typically has the form
aVariableName <- aDataReadingFunction(aPathToTheFile, optionalSettings)
- So for this to work you need the following:
 - **The file and the path to it** (Did you actually remember to download the file and do you know where it is?)
 - **The right data reading function** (What function is it that you are going to use for that particular file format? From which package does it come? Have you loaded the package (with a “library” command)? Have you installed the package at all?)
 - **Figure what optional setting to pass to the data reading function** (Are you going to skip some lines at the top? Are there multiple sheet you could read in? What are the types of the columns? Do you need to give them explicitly?)

Reading data into R - example

- Here is an example:

```
read.csv("myfile.csv", header = FALSE)
```

- This reads the "myfile.csv"
- It read it as a csv file (due to the use of the function `read.csv`)
- It has the extra argument `header = FALSE` that tells the function that the columns in our csv files does not have names that appear in the first row
 - To have such names is standard, if we only need to pass this argument if we do not have them – the **default argument** is `header = TRUE`
 - Note, that many functions have default arguments for some their parameters/arguments. These default can be seen on the help page for the function
- Instead of a file name as in "myfile.csv" you can also pass a url to download data from the web into R.

Reading data into R

- Good steps for loading in data
 - 1) Figure out what format the file is
 - 2) Figure out the right data reading function from the right package to load this file
 - 3) Have a look at the file (in Excel, LibreOffice, a text editor, ... etc)
 - 4) Determine whether there are any specific arguments you need to provide to the data reading function
 - 5) When you have loaded the data into R, have a look at it (using head, tail, str, etc.)
 - 6) Is the data as expected? If not, repeat step 3-6.

Reading data into R

- To read other data types into R, you often need specific packages
 - For instance, for excel files, you can use the function `read_excel` from the “readxl” package (See the ETL notebook of today)
- To read data bases you can use the package “RODBC”
 - First you create/open a connection to the database using the function `odbcConnect`
 - Then you can send SQL queries to the data base using the established connection and the function `sqlQuery`, which return a data frame
 - Finally, you should remember to `close` the connection using the `close` function
 - For details see the book.

Transforming data

- It can be hard and time consuming
- It might be more of an art than a science
- It can require knowledge about the data
- Having a clear idea about what tidy data is make it easier to turn data into tidy data
- We will first discuss how to get information out of data by doing other transformations using the functions `filter`, `arrange`, `select`, `mutate`, `group_by`, and `summarise`
- We will then discuss ways of making data tidy with the `gather` and `spread` functions

Transforming data

- With the dplyr package we can easily:
 - Pick observations by their value (`filter`)
 - Reorder the rows (`arrange`)
 - Pick variables by their names (`select`)
 - Create new variables with functions of existing variables (`mutate`)
 - Collapse many values down to a single summary (`summarise`), maybe per group (`group_by`)
- All these operations are extremely useful

The `filter` function

- `filter` allow us to filter out some of the rows, depending on their particular values
 - We filter out rows based on an expression
 - Expressions are statements about values of variables and combinations of such
- Examples:
 - *All flights on August 20, 2013*
 - *All flights in 2014*
 - *All flights which departed before noon*
 - *All flights with a delay of more than 2 hours*
 - *All flights on August 20, 2013 that departed before noon and were delayed by more than 2 hours.*

```
library(dplyr)
library(nycflights13)
flights
```

| year | month | day | dep_time | sched_dep_time | dep_delay | arr_time | sched_arr_time | arr_delay | carrier | flight | tailnum | origin | dest | air_time | distance | hour |
|------|-------|-----|----------|----------------|-----------|----------|----------------|-----------|---------|--------|---------|--------|------|----------|----------|------|
| 2013 | 1 | 1 | 517 | 515 | 2 | 830 | 819 | 11 | UA | 1545 | N14228 | EWB | IAH | 227 | 1400 | 5 |
| 2013 | 1 | 1 | 533 | 529 | 4 | 850 | 830 | 20 | UA | 1714 | N24211 | LGA | IAH | 227 | 1416 | 5 |
| 2013 | 1 | 1 | 542 | 540 | 2 | 923 | 850 | 33 | AA | 1141 | N619AA | JFK | MIA | 160 | 1089 | 5 |
| 2013 | 1 | 1 | 544 | 545 | -1 | 1004 | 1022 | -18 | B6 | 725 | N804JB | JFK | BQN | 183 | 1576 | 5 |

filter examples

```
filter(flights, month == 8, day == 20)
```

| year | month | day | dep_time | sched_dep_time | dep_delay | arr_time | sched_arr_time | arr_delay | carrier | flight | tailnum | origin | dest | air_time | distance | hour |
|------|-------|-----|----------|----------------|-----------|----------|----------------|-----------|---------|--------|---------|--------|------|----------|----------|------|
| 2013 | 8 | 20 | 54 | 2359 | 55 | 432 | 344 | 48 | B6 | 1503 | N615JB | JFK | SJU | 196 | 1598 | 23 |
| 2013 | 8 | 20 | 455 | 500 | -5 | 618 | 642 | -24 | US | 1993 | N155UW | EWL | CLT | 72 | 529 | 5 |
| 2013 | 8 | 20 | 526 | 530 | -4 | 753 | 801 | -8 | UA | 1545 | N24211 | EWL | IAH | 186 | 1400 | 5 |
| 2013 | 8 | 20 | 536 | 540 | -4 | 832 | 840 | -8 | AA | 701 | N5CFAA | JFK | MIA | 148 | 1089 | 5 |

```
filter(flights, arr_delay >= 120)
```

| year | month | day | dep_time | sched_dep_time | dep_delay | arr_time | sched_arr_time | arr_delay | carrier | flight | tailnum | origin | dest | air_time | distance | hour |
|------|-------|-----|----------|----------------|-----------|----------|----------------|-----------|---------|--------|---------|--------|------|----------|----------|------|
| 2013 | 1 | 1 | 811 | 630 | 101 | 1047 | 830 | 137 | MQ | 4576 | N531MQ | LGA | CLT | 118 | 544 | 6 |
| 2013 | 1 | 1 | 848 | 1835 | 853 | 1001 | 1950 | 851 | MQ | 3944 | N942MQ | JFK | BWI | 41 | 184 | 18 |
| 2013 | 1 | 1 | 957 | 733 | 144 | 1056 | 853 | 123 | UA | 856 | N534UA | EWL | BOS | 37 | 200 | 7 |
| 2013 | 1 | 1 | 1114 | 900 | 134 | 1447 | 1222 | 145 | UA | 1086 | N76502 | LGA | IAH | 248 | 1416 | 9 |

Expressions in R

- Comparisons
 - **>**, **>=**, **<**, **<=**: Greater than, greater than or equal, less than, less than or equal,
 - **==**, **!=**: Equal, not equal
 - Note: use “==” for character strings, truth values, and integers. (“hej” == “hej”, “x == TRUE”), but not for floating point numbers ($1/49 * 49 == 1$) or missing values ($x == NA$). For test of NA use `is.na(x)` instead.
- Logical operators
 - **&**: and. Ex: `filter(flights, month == 8 & day == 20)`
 - **|**: or. Ex: `filter(flights, month == 8 | day == 20)`
 - **!**: not. Ex: `filter(flights, !(arr_delay < 120))`
 - (**all**: all)
 - (**any**: any)
 - Rules: $!(x == y)$ same as $x != y$, $!(A \& B)$ same as $!A | !B$, $!(A | B)$ same as $!A \& !B$, $!(x > y)$ same as $x <= y$, $!(x <= y)$ same as $x > y$, etc.
- Missing values, use `is.na()`
 - `filter(flights, !is.na(carrier))`

The select function

- An easy way of selecting only some columns
- Can work with `starts_with()`, `ends_with()`, `contains()`

```
> table1
# A tibble: 6 × 4
  country year cases population
  <chr> <int> <int> <int>
1 Afghanistan 1999 745 19987071
2 Afghanistan 2000 2666 20595360
3 Brazil 1999 37737 172006362
4 Brazil 2000 80488 174504898
5 China 1999 212258 1272915272
6 China 2000 213766 1280428583
> select(table1, year, cases)
# A tibble: 6 × 2
  year cases
  <int> <int>
1 1999 745
2 2000 2666
3 1999 37737
4 2000 80488
5 1999 212258
6 2000 213766
> select(table1, -year)
# A tibble: 6 × 3
  country cases population
  <chr> <int> <int>
1 Afghanistan 745 19987071
2 Afghanistan 2666 20595360
3 Brazil 37737 172006362
4 Brazil 80488 174504898
5 China 212258 1272915272
6 China 213766 1280428583
```

```
> select(table1, year:population)
# A tibble: 6 × 3
  year cases population
  <int> <int> <int>
1 1999 745 19987071
2 2000 2666 20595360
3 1999 37737 172006362
4 2000 80488 174504898
5 1999 212258 1272915272
6 2000 213766 1280428583
> select(table1, starts_with("c"))
# A tibble: 6 × 2
  country cases
  <chr> <int>
1 Afghanistan 745
2 Afghanistan 2666
3 Brazil 37737
4 Brazil 80488
5 China 212258
6 China 213766
> select(table1, 1:3)
# A tibble: 6 × 3
  country year cases
  <chr> <int> <int>
1 Afghanistan 1999 745
2 Afghanistan 2000 2666
3 Brazil 1999 37737
4 Brazil 2000 80488
5 China 1999 212258
6 China 2000 213766
```

The mutate function

- Add new columns (or transforms old) based on combinations/calculations of old ones
- Transmute can be used if you only want to keep the mentioned columns

```
> testDF
  date clicks impressions
1 2016-09-11    12      74
2 2016-09-12     1      65
3 2016-09-13    18       2
4 2016-09-14    12     143
5 2016-09-15    27     152
6 2016-09-16    23     106
7 2016-09-17    19     177
8 2016-09-18    27     108
9 2016-09-19    24      46
10 2016-09-20    NA      67
11 2016-09-21    19      72

> mutate(testDF, clickThroughRate = clicks / impressions)
  date clicks impressions clickThroughRate
1 2016-09-11    12      74      0.16216216
2 2016-09-12     1      65      0.01538462
3 2016-09-13    18       2      9.00000000
4 2016-09-14    12     143      0.08391608
5 2016-09-15    27     152      0.17763158
6 2016-09-16    23     106      0.21698113
7 2016-09-17    19     177      0.10734463
8 2016-09-18    27     108      0.25000000
9 2016-09-19    24      46      0.52173913
10 2016-09-20    NA      67             NA
11 2016-09-21    19      72      0.26388889

> mutate(testDF, shareOfTotalClicks =
+         clicks / sum(clicks, na.rm = TRUE))
  date clicks impressions shareOfTotalClicks
1 2016-09-11    12      74      0.065934066
2 2016-09-12     1      65      0.005494505
3 2016-09-13    18       2      0.098901099
4 2016-09-14    12     143      0.065934066
5 2016-09-15    27     152      0.148351648
6 2016-09-16    23     106      0.126373626
7 2016-09-17    19     177      0.104395604
8 2016-09-18    27     108      0.148351648
9 2016-09-19    24      46      0.131868132
10 2016-09-20    NA      67             NA
11 2016-09-21    19      72      0.104395604

> mutate(testDF, impressions = impressions * 100)
  date clicks impressions
1 2016-09-11    12     7400
2 2016-09-12     1     6500
3 2016-09-13    18     200
4 2016-09-14    12    14300
5 2016-09-15    27    15200
6 2016-09-16    23    10600
7 2016-09-17    19    17700
8 2016-09-18    27    10800
9 2016-09-19    24     4600
10 2016-09-20    NA     6700
11 2016-09-21    19     7200
```


The `arrange` function

- A bit like filter, but instead of removing rows, it orders the rows.
- It takes column names and orders the rows by those.
- Default is to order in increasing order, but `desc` can be used to order in descending order.
- NAs are put at the end

```
> arrange(testDF, clicks, impressions)
  date clicks impressions
1 2016-09-12      1         65
2 2016-09-11     12         74
3 2016-09-14     12        143
4 2016-09-13     18          2
5 2016-09-21     19         72
6 2016-09-17     19        177
7 2016-09-16     23        106
8 2016-09-19     24          46
9 2016-09-18     27        108
10 2016-09-15     27        152
11 2016-09-20    NA          67
> arrange(testDF, desc(impressions))
  date clicks impressions
1 2016-09-17     19        177
2 2016-09-15     27        152
3 2016-09-14     12        143
4 2016-09-18     27        108
5 2016-09-16     23        106
6 2016-09-11     12          74
7 2016-09-21     19          72
8 2016-09-20    NA          67
9 2016-09-12      1          65
10 2016-09-19     24          46
11 2016-09-13     18           2
```

The summarise and group_by functions

```
> testDF
  date   clicks impressions channel
1 2016-09-11     12         74  social
2 2016-09-12      1         65 display
3 2016-09-13     18          2 display
4 2016-09-14     12        143  social
5 2016-09-15     27        152  social
6 2016-09-16     23        106 display
7 2016-09-17     19        177  social
8 2016-09-18     27        108 display
9 2016-09-19     24         46  social
10 2016-09-20     NA         67  social
11 2016-09-21     19         72  social
```

```
> testDF %>%
+   group_by(channel) %>%
+   summarise(clicks = sum(clicks, na.rm = TRUE),
+             impressions = sum(impressions),
+             count = n())
# A tibble: 2 x 4
  channel clicks impressions count
  <chr>   <int>         <int> <int>
1 display     69          281     4
2  social    113          731     7
```

```
> summarise(testDF,
+           clicksTotal = sum(clicks),
+           impressionAvg = mean(impressions))
  clicksTotal impressionAvg
1          NA             92

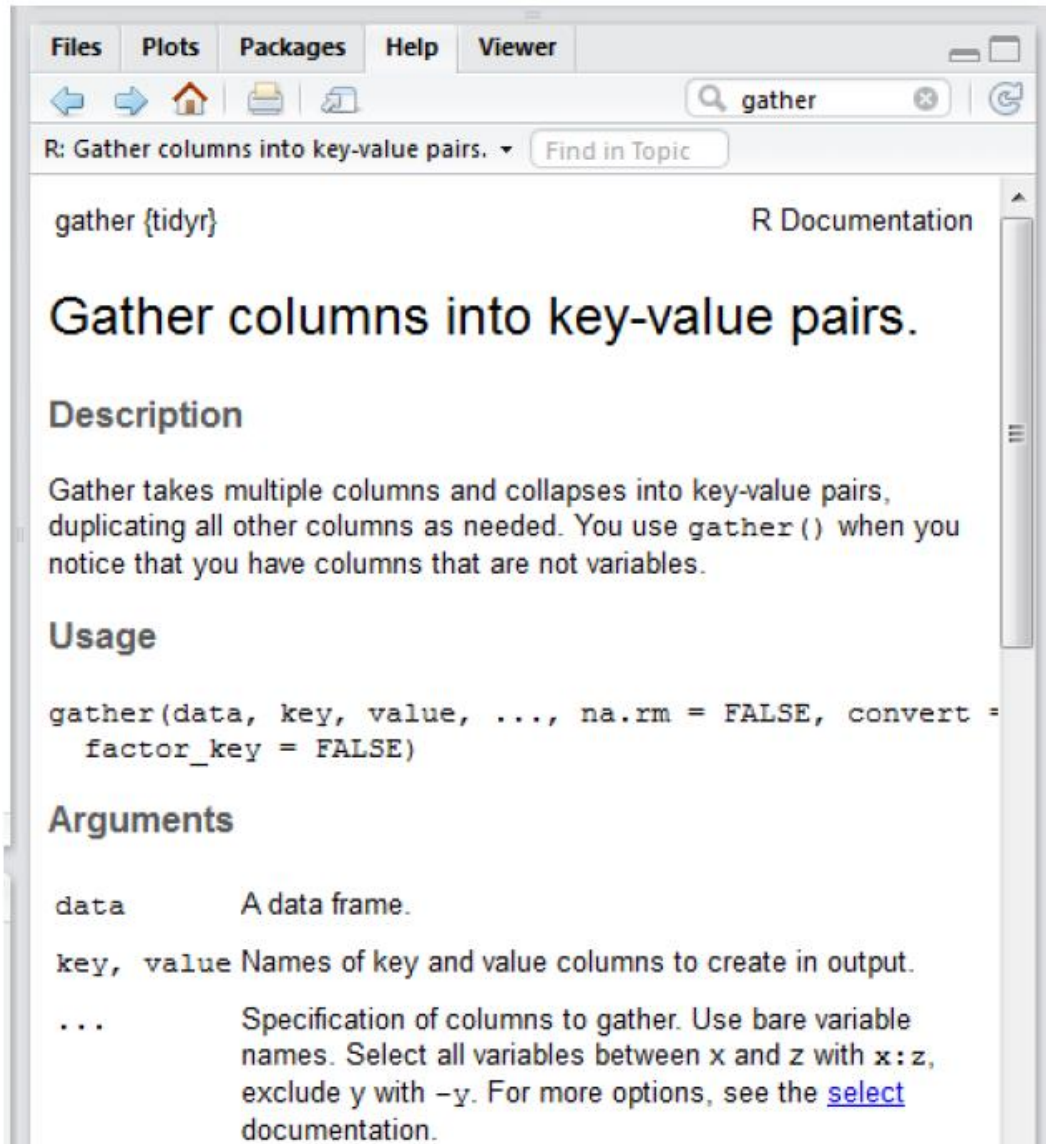
> summarise(testDF,
+           clicksTotal = sum(clicks, na.rm = TRUE),
+           impressionAvg = mean(impressions))
  clicksTotal impressionAvg
1         182             92
```

- Aggregate variables to a single value
- With `group_by`, it can be a single value per group

The `gather` function

- Comes from the “tidyr” package.
- Makes wide tables narrower and longer.
- Useful for the case when a variable is spread across columns – the values of a variable are in the column names.
- In this case, we need to “gather” those columns into a new variable.

How to use gather



The screenshot shows the R Documentation window for the `gather` function in the `tidyr` package. The window has a menu bar with 'Files', 'Plots', 'Packages', 'Help', and 'Viewer'. Below the menu bar is a search bar with the text 'gather' and a 'Find in Topic' button. The main content area is titled 'gather {tidyr}' and 'R Documentation'. It includes a description of the function, its usage, and its arguments.

Gather columns into key-value pairs.

Description

Gather takes multiple columns and collapses into key-value pairs, duplicating all other columns as needed. You use `gather()` when you notice that you have columns that are not variables.

Usage

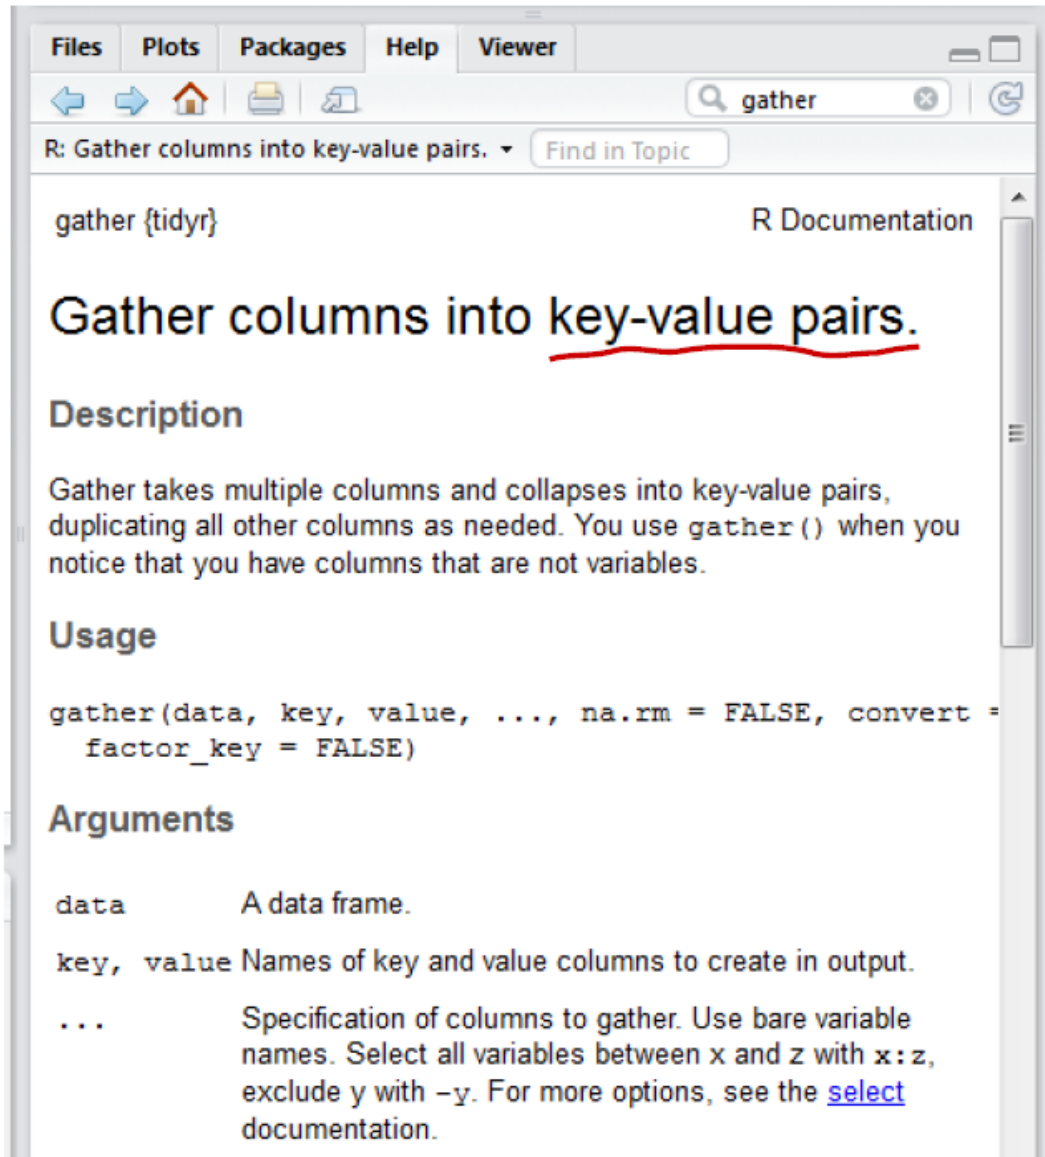
```
gather(data, key, value, ..., na.rm = FALSE, convert = FALSE, factor_key = FALSE)
```

Arguments

| | |
|---------------------------------------|---|
| <code>data</code> | A data frame. |
| <code>key</code> , <code>value</code> | Names of key and value columns to create in output. |
| <code>...</code> | Specification of columns to gather. Use bare variable names. Select all variables between <code>x</code> and <code>z</code> with <code>x:z</code> , exclude <code>y</code> with <code>-y</code> . For more options, see the select documentation. |

```
> table4a
# A tibble: 3 × 3
  country `1999` `2000`
*   <chr>   <int>  <int>
1 Afghanistan    745    2666
2      Brazil  37737   80488
3      China 212258  213766
> gather(table4a,
+         `1999`, `2000`,
+         key = "year",
+         value = "cases")
# A tibble: 6 × 3
  country  year  cases
  <chr>   <chr>  <int>
1 Afghanistan  1999    745
2      Brazil  1999  37737
3      China  1999 212258
4 Afghanistan  2000    266
5      Brazil  2000  80488
6      china  2000 213766
> |
```

How to use gather



The screenshot shows the R Documentation window for the `gather` function. The title is "gather {tidyr}" and the subtitle is "R Documentation". The main heading is "Gather columns into key-value pairs.". Below this is a "Description" section stating: "Gather takes multiple columns and collapses into key-value pairs, duplicating all other columns as needed. You use `gather()` when you notice that you have columns that are not variables." The "Usage" section shows the function signature: `gather(data, key, value, ..., na.rm = FALSE, convert = FALSE, factor_key = FALSE)`. The "Arguments" section lists: `data` (A data frame.), `key, value` (Names of key and value columns to create in output.), and `...` (Specification of columns to gather. Use bare variable names. Select all variables between `x` and `z` with `x:z`, exclude `y` with `-y`. For more options, see the [select](#) documentation.).

```
> table4a
# A tibble: 3 x 3
  country `1999` `2000`
*   <chr>   <int>   <int>
1 Afghanistan 745    2666
2   Brazil 37737  80488
3   China 212258 213766
> gather(table4a,
+         `1999`, `2000`,
+         key = "year",
+         value = "cases")
# A tibble: 6 x 3
  country year cases
  <chr>   <chr>   <int>
1 Afghanistan 1999    745
2   Brazil 1999  37737
3   China 1999 212258
4 Afghanistan 2000    2666
5   Brazil 2000  80488
6   China 2000 213766
> |
```

The spread function

- Comes from the “tidyr” package.
- Makes long tables shorter and wider.
- Useful when one observation is scattered across multiple rows.
- We need to spread the observation out

How to use spread

spread {tidyr}

R Documentation

Spread a key-value pair across multiple columns.

Description

Spread a key-value pair across multiple columns.

Usage

```
spread(data, key, value, fill = NA, convert = FALSE, drop = TRUE,
  sep = NULL)
```

Arguments

data A data frame.

key The bare (unquoted) name of the column whose values will be used as column headings.

value The bare (unquoted) name of the column whose values will populate the cells.

```
> table2
# A tibble: 12 x 4
  country year type count
  <chr> <int> <chr> <int>
1 Afghanistan 1999 cases 745
2 Afghanistan 1999 population 19987071
3 Afghanistan 2000 cases 2666
4 Afghanistan 2000 population 20595360
5 Brazil 1999 cases 37737
6 Brazil 1999 population 172006362
7 Brazil 2000 cases 80488
8 Brazil 2000 population 174504898
9 China 1999 cases 212258
10 China 1999 population 1272915272
11 China 2000 cases 213766
12 China 2000 population 1280428583
> spread(table2, key = type, value = count)
# A tibble: 6 x 4
  country year cases population
  <chr> <int> <int> <int>
1 Afghanistan 1999 745 19987071
2 Afghanistan 2000 2666 20595360
3 Brazil 1999 37737 172006362
4 Brazil 2000 80488 174504898
5 China 1999 212258 1272915272
6 China 2000 213766 1280428583
```

How to use spread

spread {tidyr}

R Documentation

Spread a key-value pair across multiple columns.

Description

Spread a key-value pair across multiple columns.

Usage

```
spread(data, key, value, fill = NA, convert = FALSE, drop = TRUE,
  sep = NULL)
```

Arguments

data A data frame.

key The bare (unquoted) name of the column whose values will be used as column headings.

value The bare (unquoted) name of the column whose values will populate the cells.

```
> table2
# A tibble: 12 x 4
  country year type count
  <chr> <int> <chr> <int>
1 Afghanistan 1999 cases 745
2 Afghanistan 1999 population 19987071
3 Afghanistan 2000 cases 2666
4 Afghanistan 2000 population 20595360
5 Brazil 1999 cases 37737
6 Brazil 1999 population 172006362
7 Brazil 2000 cases 80488
8 Brazil 2000 population 174504898
9 China 1999 cases 212258
10 China 1999 population 1272915272
11 China 2000 cases 213766
12 China 2000 population 1280428583
```

```
> spread(table2, key = type, value = count)
# A tibble: 6 x 4
  country year cases population
  <chr> <int> <int> <int>
1 Afghanistan 1999 745 19987071
2 Afghanistan 2000 2666 20595360
3 Brazil 1999 37737 172006362
4 Brazil 2000 80488 174504898
5 China 1999 212258 1272915272
6 China 2000 213766 1280428583
```

Writing data from R to files

- CSV is a very common format so it is often the best way to write a data frame to a file

- Example:

```
write.csv(report, "revenue_report.csv", row.names = FALSE)
```

- This writes out the data frame report
- It writes it in csv format to the file revenue_report.csv
- It makes sure there will not be added a column with row numbers (or row names if the data frame have such

More literature on data transformation

- “Introduction to R for Business Intelligence” by Jay Gendron
 - Chapters 1 & 2
- “Tidy data” by Hadley Wickham
 - <https://www.jstatsoft.org/index.php/jss/article/view/v059i10/v59i10.pdf>
- “R for Data Science” by Garrett Grolemund and Hadley Wickham
 - Chapters 5, 9-16
 - <http://r4ds.had.co.nz/>
- “An Introduction to Data Science (version 3)” by Jeffrey Stanton
 - Chapters 1 & 5
 - <https://drive.google.com/file/d/0B6iefdnF22XQeVZDSkxjZ0Z5VUE/edit>

Exercise

- Open the notebook “ETL example and exercise notebook” from the BIBA-2018 Library (you might need to clone it again to get the new notebook)
- Read and run the notebook cells with code one by one.
- Do the exercises in the notebook as well

Data cleaning

Cleaning data

- Together with data transformation this is the time consuming tasks that always needs to be done, but most business people forget to account for
- Cleaning data amounts to the following at least
 - Summarizing your data for inspection
 - Finding and fixing flawed data
 - Converting inputs to data types suitable for analysis
 - Adapting string variables to a standard
- The author of the text book operates with a structured process he calls ***summarize-fix-convert-adapt***
- Look at the notebook “Data cleaning notebook” from the BIBA-2018 Library on Azure Notebooks: <https://notebooks.azure.com/jensuh/libraries/BIBA-2018/html/Data%20cleaning%20notebook.ipynb>

Exercise

- Find a data set that could be useful for your business case and load it into R
- Do some simple investigation, transformation and cleaning of the data